


European
Commission

Horizon 2020
European Union funding
for Research & Innovation

Introduction to Optimal Control



Ludovic Righetti (NYU)
Nicolas Mansard (CNRS)



- ❑ What can we do with optimal control?
- ❑ Where is optimal-control in the robot galaxy?
- ❑ What is dynamic programming?
- ❑ Should you shoot or collocate?
- ❑ Why make your dynamic program *differential*?
- ❑ Is multiple shooting about guns?
- ❑ What is Crocoddyl good for, and what is beyond?



What can we do with optimal control?

VIDEO INTRODUCTION



Information Theoretic Model
Predictive Control
[Williams et al. 2018]





OC with Linear Inverted
Pendulum Model
[Herdt et al. 2010]



OC with Centroidal Momentum Dynamics and
Full Body Kinematics
[Ponton et al. 2018], [Carpentier et al. 2018],
[Dai et al. 2014], [Herzog et al. 2015]





Full-body Optimal Control

Synthesis and stabilization of complex behaviors with online trajectory optimization

Yuval Tassa, Tom Erez and Emo Todorov

Movement Control Laboratory
University of Washington

IROS 2012

[Tassa et al. 2010]
DDP with Full-Body Dynamics
(realtime control)

[Mordatch et al. 2012]
Nonlinear Optimization for
Multi-Contact Tasks

Discovery of complex behaviors through Contact-Invariant Optimization

Igor Mordatch, Emo Todorov and Zoran Popovic

Movement Control Laboratory and GRAIL
University of Washington

SIGGRAPH 2012



What is dynamic programming

INTRODUCTION TO BELMAN'S EQUATIONS

Optimal control problem

$$\min_{\underbrace{u_0, \dots, u_{T-1}}_{\text{Find control inputs to minimize cost}}} \sum_{t=0}^{T-1} \underbrace{l_t(x_t, u_t)}_{\text{stage costs}} + \underbrace{l_T(x_T)}_{\text{terminal cost}}$$


$$x_{t+1} = f_t(x_t, u_t)$$

deterministic dynamics

$$g(x_t, u_t) \leq 0$$


state and control constraints






$$\{x\} = x_0, \dots, x_T$$

$$\{u\} = u_0, \dots, u_{T-1}$$




$$\{x\} = x_0, \dots, x_T$$

$$\{u\} = u_0, \dots, u_{T-1}$$



$$\{x'\} = x'_0, \dots, x'_T$$
$$\{u'\} = u'_0, \dots, u'_{T-1}$$





$\pi(x)$ \Rightarrow control policy


$$\{u\}^*$$

the optimal control trajectory

$$\pi^*(x)$$

the optimal control policy






Principle of Optimality

How can we find the optimal control?

The Principle of Optimality breaks down the problem



Subpath of optimal paths are also optimal for their own subproblem





Principle of Optimality

How can we find the optimal control?


The Principle of Optimality breaks down the problem

Optimal Cost to
Go or Value
Function

$$V_t(x_t) = \min_{u_t, \dots, u_{N-1}} \sum_{k=t}^{T-1} l_k(x_k, u_k) + l_T(x_T)$$


Bellman's
Principle of
Optimality

$$V_t(x_t) = \min_{u_t} l_t(x_t, u_t) + V_{t+1}(x_{t+1})$$


$$x_{t+1} = f_t(x_t, u_t)$$




$$V_t(x_t) = \min_{u_t} l_t(x_t, u_t) + V_{t+1}(x_{t+1})$$



Final States
Stage T
 $V_T(x_T)$



$$V_t(x_t) = \min_{u_t} l_t(x_t, u_t) + V_{t+1}(x_{t+1})$$




Stage T-1 Final States
Stage T

$$V_{T-1}(x_{T-1}) \quad V_T(x_T)$$


$$\pi_{T-1}(x_{T-1})$$



$$V_t(x_t) = \min_{u_t} l_t(x_t, u_t) + V_{t+1}(x_{t+1})$$



$$V_t(x_t) = \min_{u_t} l_t(x_t, u_t) + V_{t+1}(x_{t+1})$$



Problems with linear dynamics and quadratic costs can be solved explicitly!

$$x_{t+1} = F_x x_t + F_u u_t$$

Linear dynamics

$$\min \sum_{t=0}^{T-1} (x_t^T L_x x_t + u_t^T L_u u_t) + x_T^T L_x x_T$$

Quadratic cost

$$L_x \geq 0 \quad L_u > 0$$

- Set $W_T = L_x$
- For t from $T-1$ to 0 , do backward recursion

$$\begin{aligned} K_t &= -(F_u^T W_{t+1} F_u + L_u)^{-1} F_u^T W_{t+1} F_x \\ Q_t &= L_x + F_x^T W_t + F_x + F_x^T W_{t+1} F_u K_t \end{aligned}$$

Discrete-time
Riccati equation

- The cost-to-go at stage t is $V_t(x_t) = x_t^T W_t x_t$
- The optimal policy is $\pi^*(x_t) = K_t x_t$

The policy is a linear feedback controller with gain K_t

$$\text{Bellman Equation } V_t(x_t) = \min_{u_t} l_t(x_t, u_t) + V_{t+1}(x_{t+1})$$

Problems:

- Curse of dimensionality
- minimization in Bellman equation

⇒ Approximate solution to Bellman equation
(DDP, trajectory optimization, reinforcement learning, etc)



Solving Bellman's Equations

- [1] Bonnali'19 ArX:1903.00155
- [2] Mordach'14 DOI:2185520.2185539
- [3] Posa'14 DOI:0278364913506757
- [4] Winkler'18 IEEE:2798285
- [5] Rajamaki'17 DOI:3099564.3099579

Resolution Method:
Stochastic - Deterministic

$$\text{Bellman's Equation } V_t = \min_{u_t} l(x_t, u_t) + V_{t+1}(f(x_t, u_t))$$

LQR
(exact solution)

Non LQR
(approximate solution)

Indirect Methods
Pontryagin's Maximum Principle
Rockets, Cars (small dimensions)

GuSTO [1]

Direct Methods
(Most popular in robotics)

"local"
Trajectory optimization

"global"
Value/Policy optimization

Collocation

CIO [2] TOWR [4]

TrajOpt

"Direct" trajectory optim [3]

Shooting

DDP

Multiple shooting

CMAES, PI²

Guided policy
search

Explicit MPC

Q learning


Actor Critic

DDPG, TRPO, PPO



Should we collocate or shoot?

TRANSCRIPTION



Transcribing: “representing” the reality

$$\begin{aligned} \min_{\substack{\underline{x}: t \rightarrow x(t) \\ \underline{u}: t \rightarrow u(t)}} \int_0^T l(x(t), u(t)) dt + l_T(x(T)) \\ \text{s.t. } \forall t, \dot{x}(t) = f(x(t), u(t)) \end{aligned}$$

Optimal control problem (OCP)
with continuous variables
(infinite-dimension)

$$\begin{aligned} \min_{\substack{\underline{x} = \theta_{x1} \dots \theta_{xn} \\ \underline{u} = \theta_{u1} \dots \theta_{un}}} \sum_t l(x(t|\theta), u(t|\theta)) + l_T(x(T|\theta)) \\ \text{s.t. at some } t, \dot{x}(t|\theta) = f(t|\theta_x, \theta_u) \end{aligned}$$

Nonlinear optimization problem (NLP)
with static variables
(finite dimension)

$\theta_x \theta_u$ represents the continuous $\underline{x}, \underline{u}$ trajectories




Transcription: shooting versus collocation

\underline{u} is easy to represent (piecewise polynomials)

– what about \underline{x} ?

- Collocation: \underline{x} is represented by another polynomials



Polynomials(θ_u)

Polynomials(θ_x)





Transcription: shooting versus collocation

u is easy to represent (piecewise polynomials)

– what about x ?

- ❑ Collocation: x is represented by another polynomials



Problems:

The solution to $\dot{x}(t) = f(x(t), u(t))$ is **not polynomial**

The dynamics is only checked **at some remote points**




Transcription: shooting versus collocation

\underline{u} is easy to represent (piecewise polynomials)

– what about \underline{x} ?

- ❑ Shooting: \underline{x} is represented by an integrator and only evaluated sparsely



Polynomials(θ_u)

$\theta_x = (x_I, \dots, x_T)$





Transcription: shooting versus collocation

\underline{u} is easy to represent (piecewise polynomials)

– what about \underline{x} ?

- ❑ Shooting: \underline{x} is represented by an integrator and only evaluated sparsely



Problems:

The state is **sparsely** and **approximately** known


You may need an **accurate integrator** (complex+costly)



$$\min_{\underline{u}=(u_0..u_{T-1})} \sum_t l(x(u_0..u_{t-1}|x_0), u_t) + l_T(x(u_0..u_{T-1}))$$

where $x(u_0..u_{t-1}|x_0)$ is a function of \underline{u}


- **Unconstrained** optimization
- The function $\underline{u}(\underline{x})$ is numerically **unstable**



- ❑ Easy to implement
 - ❑ Integrator, derivatives, Newton-descent
- ❑ Side effect: you can focus on efficiency
- ❑ Numerically unstable
- ❑ The initial-guess θ_{xu} should be meaningful
- ❑ At the end, maybe we don't care so much ...

Why make your dynamic program *differential*?


D.D.P.



Tassa et al., IROS' 12




1. DDP as iterative LQR



- “Next-step” is a nonlinear function

$$\Delta x' = f(x + \Delta x, u + \Delta u) - f(x, u)$$



- “Next-step” is a nonlinear function

$$\Delta x' = f(x + \Delta x, u + \Delta u) - f(x, u)$$

- Approximate by

$$\Delta x' = f(x, u) + F_x \Delta x + F_u \Delta u - f(x, u)$$

□ Nonlinear optimal control problem

$$\begin{aligned} \min_{\{x\}, \{u\}} & \sum_{t=0}^{T-1} l(x_t, u_t) + l_T(x_T) \\ \text{s.t. } & \forall t=0..T-1 \quad x_{t+1} = f(x_t, u_t) \end{aligned}$$

□ Linear-Quadratic problem ... solved in Part 1.

$$\begin{aligned} \min_{\{\Delta x\}, \{\Delta u\}} & \sum_{t=0}^{T-1} \begin{pmatrix} L_x \\ L_u \end{pmatrix}^T \begin{pmatrix} \Delta x_t \\ \Delta u_t \end{pmatrix} + \frac{1}{2} \begin{pmatrix} \Delta x_t \\ \Delta u_t \end{pmatrix}^T \begin{pmatrix} L_{xx} & L_{xu} \\ L_{ux} & L_{uu} \end{pmatrix} \begin{pmatrix} \Delta x_t \\ \Delta u_t \end{pmatrix} + \dots \\ \text{s.t. } & \forall t=0..T-1 \quad \Delta x_{t+1} = F_x \Delta x_t + F_u \Delta u_t \end{aligned}$$

□ Algorithm iLQR

Initialize with a given trajectory $\{\mathbf{x}_0\}, \{\mathbf{u}_0\}$

Repeat

 Linearize/Quadratize the OCP

 Compute the LQR policy

 Simulate (roll-out) with LQR regulator

Until local minimum is reached



2. DDP as a 2-pass algorithm

$$V_t = \min_{u_t} l(x_t, u_t) + V_{t+1}(f(x_t, u_t))$$

- Backward propagation

$$Q_t = l(x_t, u_t) + V_{t+1}(f(x_t, u_t))$$

- Greedy optimization

$$V_t = \min_{u_t} Q_t(x_t, u_t)$$



DDP as a 2-pass algorithm

$$Q = l + V'$$

$$V = \min_u Q$$





DDP as a 2-pass algorithm

- Pass 1: back-propagate an approximation of V
 - We can solve Bellman for quadratic cost and linear dynamics
- Pass 2: forward propagate gains and trajectory





DDP as a 2-pass algorithm

- Pass 1: backpropagate an approximation of V





DDP as a 2-pass algorithm

- Pass 2: forward propagate gains and trajectory



- ❑ Globalization (because nonconvexity)
- ❑ Line search
 - ❑ $u = u^* + k + K (x - x^*)$
 - ❑ $x' = f(x, u)$
- ❑ Regularization
 - ❑ $Q_{uu} = L_{uu} + F_u^T V_{xx} F_u$
 - ❑ $k = Q_{uu}^{-1} Q_u$
 - ❑ $K = Q_{uu}^{-1} Q_{ux}$



3. DDP as sparse SQP

$$\min_{\{x\}, \{u\}} \sum_{t=0}^{T-1} l(x_t, u_t) + l_T(x_T)$$

$$\text{s.t. } \forall t = 0..T-1 \quad x_{t+1} = f(x_t, u_t)$$

- Reminder
- Non linear problem

$$\begin{aligned} \min_y & l(y) \\ \text{s.t. } & f(y)=0 \end{aligned}$$

- Resulting “linearization”

$$\begin{aligned} \min_{\Delta y} & l(y) + L_y \Delta y + \frac{1}{2} \Delta y^T L_{yy} \Delta y \\ \text{s.t. } & f(y) + F_y \Delta y = 0 \end{aligned}$$

$$\begin{aligned} \min_{\Delta y} \quad & l(y) + L_y \Delta y + \frac{1}{2} \Delta y^T L_{yy} \Delta y \\ \text{s.t.} \quad & f(y) + F_y \Delta y = 0 \end{aligned}$$

□ Lagrangian on the NLP

$$\mathcal{L}(y, \lambda) = l(y) + \lambda^T f(y)$$

lagrangian

Primal variable

Dual variable (multipliers)

$$\begin{aligned} \min_{\Delta y} \quad & l(y) + L_y \Delta y + \frac{1}{2} \Delta y^T L_{yy} \Delta y \\ \text{s.t.} \quad & f(y) + F_y \Delta y = 0 \end{aligned}$$

□ Lagrangian on the QP

$$\begin{aligned} \mathcal{L}(\Delta y, \lambda) = & L_y \Delta y + \frac{1}{2} \Delta y^T L_{yy} \Delta y \\ & + \lambda^T (F_y \Delta y - f(y)) \end{aligned}$$

$$\begin{aligned} \min_{\Delta y} \quad & l(y) + L_y \Delta y + \frac{1}{2} \Delta y^T L_{yy} \Delta y \\ \text{s.t.} \quad & f(y) + F_y \Delta y = 0 \end{aligned}$$

□ Lagrangian on the QP

$$\begin{aligned} \mathcal{L}(\Delta y, \lambda) = & L_y \Delta y + \frac{1}{2} \Delta y^T L_{yy} \Delta y \\ & + \lambda^T (F_y \Delta y - f(y)) \end{aligned}$$

□ Newton step

$$\begin{pmatrix} L_{yy} & F_y^T \\ F_y & 0 \end{pmatrix} \begin{pmatrix} \Delta y \\ \lambda \end{pmatrix} = \begin{pmatrix} -L_y \\ -f(y) \end{pmatrix}$$

$$\min_{\{\Delta x\}, \{\Delta u\}} \sum_{t=0}^{T-1} \begin{pmatrix} L_x \\ L_u \end{pmatrix}^T \begin{pmatrix} \Delta x_t \\ \Delta u_t \end{pmatrix} + \frac{1}{2} \begin{pmatrix} \Delta x_t \\ \Delta u_t \end{pmatrix}^T \begin{pmatrix} L_{xx} & L_{xu} \\ L_{ux} & L_{uu} \end{pmatrix} \begin{pmatrix} \Delta x_t \\ \Delta u_t \end{pmatrix} + \dots$$

$$\text{s.t. } \forall t=0..T-1 \quad \Delta x_{t+1} = F_x \Delta x_t + F_u \Delta u_t + f_t$$

$$\begin{pmatrix} L_{yy} & F_y^T \\ F_y & 0 \end{pmatrix} \begin{pmatrix} \Delta y \\ \lambda \end{pmatrix} = \begin{pmatrix} -L_y \\ -f(y) \end{pmatrix}$$

$$\min_{\{\Delta x\}, \{\Delta u\}} \sum_{t=0}^{T-1} \begin{pmatrix} L_x \\ L_u \end{pmatrix}^T \begin{pmatrix} \Delta x_t \\ \Delta u_t \end{pmatrix} + \frac{1}{2} \begin{pmatrix} \Delta x_t \\ \Delta u_t \end{pmatrix}^T \begin{pmatrix} L_{xx} & L_{xu} \\ L_{ux} & L_{uu} \end{pmatrix} \begin{pmatrix} \Delta x_t \\ \Delta u_t \end{pmatrix} + \dots$$

$$\text{s.t.} \quad \forall t=0..T-1 \quad \Delta x_{t+1} = F_x \Delta x_t + F_u \Delta u_t + f_t$$

$$\begin{bmatrix} L_{xx} & & & L_{xu} & & & -I & F_x^T \\ & \ddots & & & \ddots & & & \ddots & \ddots \\ & & L_{xx} & & & L_{xu} & & -I & F_x^T \\ & & & L_{xx} & & & & & -I \\ L_{ux} & & & L_{uu} & & & F_u^T & & \\ & \ddots & & & \ddots & & & \ddots & \\ & & L_{ux} & & & L_{uu} & & & F_u^T \\ -I & & & & & & & & \\ F_x & -I & & F_u & & & & & \\ & \ddots & \ddots & & \ddots & & & & \\ & & F_x & -I & & F_u & & & \end{bmatrix} \begin{bmatrix} \Delta x_0 \\ \vdots \\ \Delta x_{T-1} \\ \Delta x_T \\ \Delta u_0 \\ \vdots \\ \Delta u_{T-1} \\ \lambda_0 \\ \lambda_1 \\ \vdots \\ \lambda_{T-1} \end{bmatrix} = - \begin{bmatrix} L_x \\ \vdots \\ L_x \\ L_x \\ L_u \\ \vdots \\ L_u \\ f_0 \\ f_1 \\ \vdots \\ f_{T-1} \end{bmatrix}$$


$$\begin{bmatrix} L_{xx} & & & L_{xu} & & & -I & F_x^T \\ & \ddots & & & \ddots & & & \ddots \\ & & L_{xx} & & & L_{xu} & & -I & F_x^T \\ & & & L_{xx} & & & & & -I \\ L_{ux} & & & L_{uu} & & & F_u^T & & \\ & \ddots & & & \ddots & & & \ddots & \\ & & L_{ux} & & & L_{uu} & & & F_u^T \\ -I & & & & & & & & \\ F_x & -I & & F_u & & & & & \\ & \ddots & \ddots & & \ddots & & & & \\ & & F_x & -I & & F_u & & & \end{bmatrix} \begin{bmatrix} \Delta x_0 \\ \vdots \\ \Delta x_{T-1} \\ \Delta x_T \\ \Delta u_0 \\ \vdots \\ \Delta u_{T-1} \\ \lambda_0 \\ \lambda_1 \\ \vdots \\ \lambda_{T-1} \end{bmatrix} = - \begin{bmatrix} L_x \\ \vdots \\ L_x \\ L_x \\ L_u \\ \vdots \\ L_u \\ f_0 \\ f_1 \\ \vdots \\ f_{T-1} \end{bmatrix}$$

What is Crocoddyl good for, and what is beyond?


BEYOND DDP



Multiple shooting




Single shooting
“Your control is bad! “




Multiple shooting
“Your control is bad! “


Multiple shooting



“Your control is bad! “



“Still bad, but better“




Polynomials(θ_u)

Polynomials(θ_x)

❑ Collocation:

We have state and control trajectories
... and they do not match





Polynomials(θ_u)


Polynomials(θ_x)

❑ Shooting:

We have a set of state points
... and the integrator does not reach them



Example of jumping



Single

Multiple

Thanks Rohan for the illustration



$$\begin{aligned} \min_{\{x\}, \{u\}} & \sum_{t=0}^{T-1} l(x_t, u_t) + l_T(x_T) \\ \text{s.t. } & \forall t=0..T-1 \quad x_{t+1} = f(x_t, u_t) \\ & \forall t=0..T \quad g(x_t, u_t) \leq 0 \end{aligned}$$

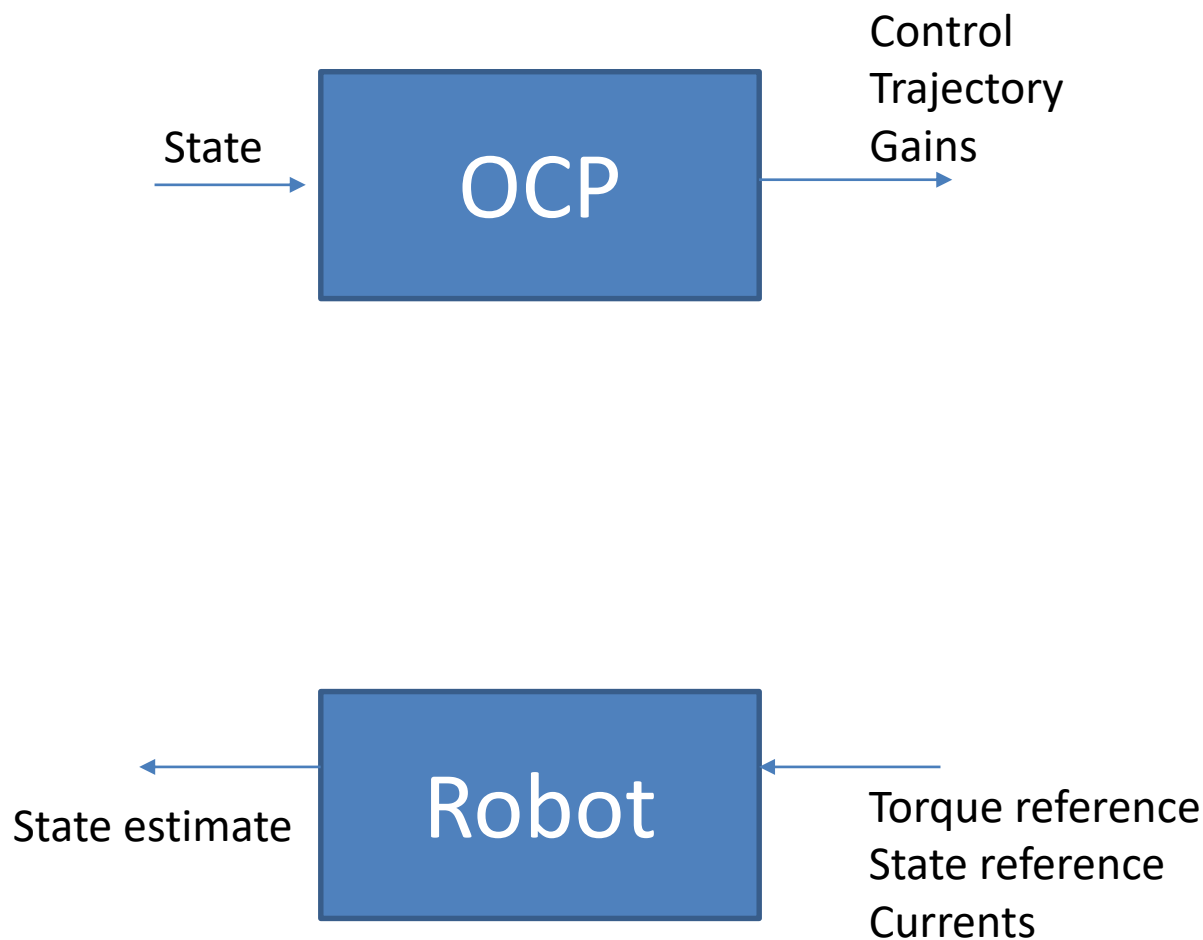
By projection

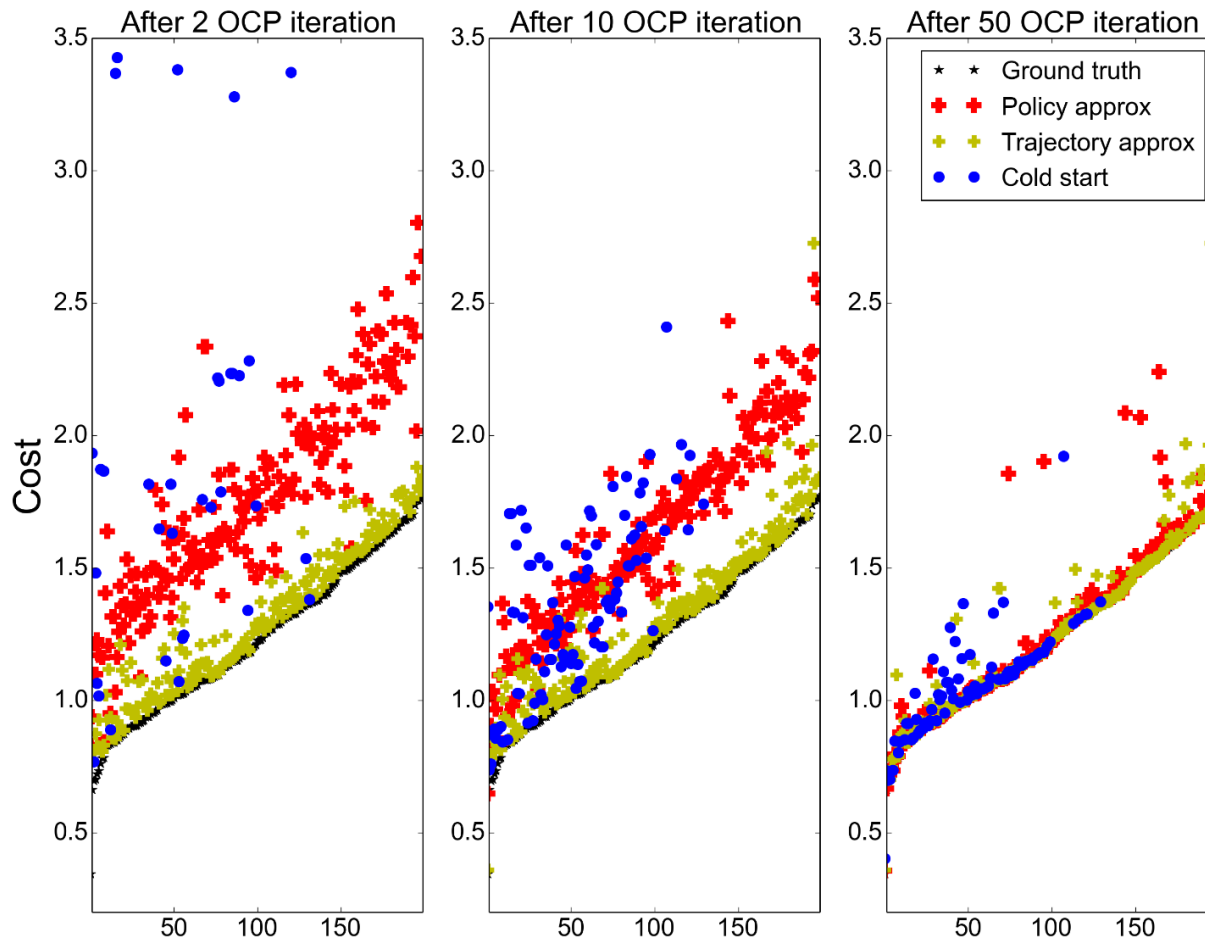
SQP, active set

By penalty

Interior point, augmented
lagrangian

❑ Closing the loop on the real robot







Start to warm-up your fingers

THE END




Numerical problems (few/none discrete constraints)

- nonconvex ... warm start needed
- very constrained ... mostly feasibility problems

The formulation/transcription is our central problem

- expert+math knowledge
- keep generalization



Optimal control = reinforcement learning

- train offline
- generalize online

